

# **Bulk Ingestion Services**

Jared Gray

29 October 2019

Draft 0.6

# Contents

<b>1. Introduction</b>	<b>1</b>
<b>2. WAN Ingestion Service</b>	<b>1</b>
2.1. Creating a New WAN Ingestion Job .....	2
2.2. Executing a WAN Ingestion Job .....	3
2.3. Describing a Single WAN Ingestion Job .....	3
2.4. Listing WAN Ingestion Jobs .....	4
2.5. Getting Data for a Chunk .....	5
<b>3. LAN Ingestion Service</b>	<b>6</b>
3.1. Creating a New LAN Ingestion Job .....	6
3.2. Executing a LAN Ingestion Job .....	7
3.3. Describing a Single LAN Ingestion Job .....	7
3.4. Listing LAN Ingestion Jobs .....	8
<b>4. Use Case: 1 TB, Two Data Centers</b>	<b>9</b>



## 1. Introduction

We define an API herein for the bulk ingestion of data over a network. Our conceptual model involves a single external data source and a potentially large number of local sinks. It is assumed that copying data from the external source is slow but that throughput between local nodes is fast.

According to this API, the ingestion proceeds in two discrete phases. In the first phase, a (proportionally small) number of nodes copy the data from the remote source and ensure their copies are consistent with one another. During the second phase, we broadcast the data obtained from the previous phase to a (proportionally large) number of nodes within the local geographic site.

## 2. WAN Ingestion Service

The *Wide Area Network (WAN) Ingestion Service* copies all data from a remote source onto the local storage of several on-site nodes. Ingested data are broken up into a series of labeled **chunks** for easier consumption by downstream services. Time-to-live mechanics are provided to free up disk space for future data.

Every data node that implements this service maintains a full copy of the remote data on local disk. It is assumed that the remote data source may be mutable. Supposing a mutation occurs in the data source following an ingestion process, the *WAN Ingestion Service* will not reflect the mutation.

Within the *WAN Ingestion Service*, all constituent service nodes maintain identical chunk data. If the data source is mutated during the ingestion process, the copies we store are consistent with one another (but not necessarily with the source). When we increase capacity for the *WAN Ingestion Service*, the new data nodes(s) must successfully ingest all live chunks before becoming available.

### 2.1. Creating a New WAN Ingestion Job

With the CREATEJOB API, a user defines a new WAN ingestion job. Its input parameters are as follows:

Name	Type	Description	Required?
sourceUri	String	URI for the data source	Yes
chunkSizeUnit	Enum	One of {KB, MB, GB, TB}	No
chunkSizeValue	Integer	Magnitude of the chunk size	No
description	String	Human-readable characterization of the job	No
timeToLiveUnit	Enum	One of {SECOND, MINUTE, HOUR, DAY}	No
timeToLiveValue	Integer	Magnitude of the time-to-live	No

Table 1: Request Parameters for CREATEJOB

In the event of successful request creation (HTTP status code 201), the API returns an identifier for the job (`jobId`) that was created. This API may return a status code of 400 when input validation fails.

## 2.2. Executing a WAN Ingestion Job

We define the EXECUTEJOB API to allow the user to execute an existing ingestion job. Its sole input parameter is a `jobId`. The output parameters are as follows:

Name	Type	Description	Required?
status	Enum	Execution status: {SUCCESS, FAILURE}	Yes
createdChunks	List [Chunk]	Chunks successfully created and replicated	Yes
failedChunks	List [Chunk]	Chunks created but unsuccessfully replicated	No

Table 2: Response Parameters for EXECUTEJOB

Accordingly, we define the Chunk API shape as follows:

Name	Type	Description	Required?
chunkId	String	Identifies the chunk	Yes
sizeInBytes	Integer	Number of bytes within the chunk payload	Yes

Table 3: Type Definition for Chunk

When the job’s execution fails, the HTTP status reflects this as a 5XX error code. Otherwise, the result’s status code is 200.

## 2.3. Describing a Single WAN Ingestion Job

The DESCRIBEJOB API takes a `jobId` as input and outputs a description of the job as follows:

Name	Type	Description	Required?
sourceURI	String	Uniform resource indicator for the data source	Yes

Name	Type	Description	Required?
chunkSize-InBytes	Integer	Maximum chunk size in bytes	Yes
numChunks	Integer	Total number of chunks in the job	No
description	String	Human-readable description of the job	No
creationDate	Date	When the job was created	Yes
timeToLiveInSeconds	Integer	Number of seconds after which this job is eligible for garbage collection	No
status	Enum	Most recent job execution status: {SUCCESS, FAILURE}	No
replicatedChunks	List [Chunk]	Successfully replicated chunks	No

Table 4: Response Parameters for DESCRIBEJOB

When a job with the specified id does not exist, the HTTP status code in the result is 404. Otherwise, the status code is 200.

## 2.4. Listing WAN Ingestion Jobs

LISTJOBS allows a user to page through all jobs that have been created and/or executed in the system.

It takes a single optional input parameter: **pageToken**. When specified, **pageToken** is used to provide the next page of results. Only recent page tokens returned in previous queries are considered valid.

The output of LISTJOBS consists of the parameters below. This API should always return an HTTP status code of 200 unless the page token is invalid (code 400).

Name	Type	Description	Required?
jobs	List [ListedJob]	WAN ingestion jobs	Yes
pageToken	String	Token for next page of results	No

Table 5: Response Parameters for LISTJOBS

We define the ListedJob API shape as follows:

Name	Type	Description	Required?
jobId	String	Identifies the WAN ingestion job	Yes
executedSuccessfully	Boolean	True if and only if the job has been executed successfully	Yes

Table 6: Type Definition for ListedJob

## 2.5. Getting Data for a Chunk

To fetch the data within a chunk, provide the `chunkId` to `GETCHUNK`. The output of `GETCHUNK` is defined below. When the specified chunk does not exist, the HTTP status code returned is 404. Otherwise, the status code shall be 200.

Name	Type	Description	Required?
chunkDataAsBase64	String	Payload of the chunk, encoded with Base64	Yes
chunkSizeInBytes	Integer	Number of bytes within this chunk	Yes

Table 7: Response Parameters for GETCHUNK



### 3. LAN Ingestion Service

The *Local Area Network (LAN) Ingestion Service* replicates data chunks gleaned from the *WAN Ingestion Service* to a series of nodes within a local data center, rack, or other co-geolocated deployment. This ingestion phase proceeds in a series of successive execution “waves.”

Like deployments in a continuous delivery environment, users can define a threshold at which a replication wave is considered successful. Failed executions can be retried by re-executing the failed job.

#### 3.1. Creating a New LAN Ingestion Job

We define the CREATEJOB API to allow a user to create a new ingestion request. Its input parameters are as follows:

Name	Type	Description	Required?
parentJobId	String	Associated WAN ingestion job	Yes
batchSize	Integer	Maximum number of hosts in each replication wave	Yes
successThreshold	Float	Fraction of hosts whose replication must succeed in each wave (a number between 0 and 1)	Yes
description	String	Human-readable description of the job	No

Table 8: Request Parameters for CREATEJOB

In the event of successful request creation (HTTP status code 201), the API returns an identifier for the job (jobId) that

was created. This API may return a status code of 400 when input validation fails.

### 3.2. Executing a LAN Ingestion Job

EXECUTEJOB plans a series of successive ingestion “waves” consistent with the job parameters then executes each wave serially. Its only input parameter is `jobId`, which is required.

The output parameters for EXECUTEJOB are as follows:

Name	Type	Description	Required?
status	Enum	Job execution status: {SUCCESS, FAILURE}	Yes
num-FailedBatches	Integer	Number of batches that failed to execute successfully	No

Table 9: Response Parameters for EXECUTEJOB

When the job’s execution fails, the HTTP status reflects this as a 5XX error code. Otherwise, the result’s status code is 200.

### 3.3. Describing a Single LAN Ingestion Job

The DESCRIBEJOB API takes a `jobId` as input and outputs a description of the job as follows:

Name	Type	Description	Required?
jobId	String	Identifies this job	Yes
parentJobId	String	Associated WAN ingestion job	Yes
batchSize	Integer	Max. number of hosts in each replication wave	Yes

Name	Type	Description	Required?
successThreshold	Float	Min. percentage of hosts whose replication must succeed in each wave	Yes
creationDate	Date	When the job was created	Yes
numExecutions	Integer	Number of attempted job executions	Yes
status	Enum	Most recent execution status: {SUCCESS, FAILURE}	No

Table 10: Response Parameters for DESCRIBEJOB

When a job with the specified id does not exist, the HTTP status code in the result is 404. Otherwise, the status code is 200.

### 3.4. Listing LAN Ingestion Jobs

LISTJOBS allows a user to page through all jobs that have been created and/or executed in the system.

It takes a single optional input parameter: **pageToken**. When specified, **pageToken** is used to provide the next page of results. Only recent page tokens returned in previous queries are considered valid.

The output of LISTJOBS consists of the parameters below. This API should always return an HTTP status code of 200 unless the page token is invalid (400).

Name	Type	Description	Required?
jobs	List [ListedJob]	LAN ingestion jobs	Yes
pageToken	String	Token to fetch next page of results	No

Table 11: Response Parameters for LISTJOBS

We define the `ListedJob` API shape as follows:

Name	Type	Description	Required?
<code>jobId</code>	String	Identifies the LAN ingestion job	Yes
<code>executedSuccessfully</code>	Boolean	True if and only if the job has been executed successfully.	Yes

Table 12: Type Definition for `ListedJob`

## 4. Use Case: 1 TB, Two Data Centers

The primary use case whose requirements drove the initial API for these services involves two data centers with 50 thousand nodes in each. To facilitate this use case, we would deploy both the *LAN Ingestion Service* and the *WAN Ingestion Service* in each data center. There are at least two different ways to proceed from here.

One choice we could make is to create two WAN ingestion jobs and execute them concurrently in each data center. Alternatively, we could wait for one WAN ingestion job to finish in the first data center, then use DC1 as the source for a second WAN ingestion to the other data center. Which choice we make will depend primarily on the networking characteristics between data centers versus the external data source. Either way, we can then execute LAN ingestion jobs at both sites to complete replication.

This design strives to generalize. It is our goal to support data ingestion processes with an arbitrary number of data

centers and/or data sources. We also hope to provide flexibility for fluctuating network conditions and future requirements.